

Polynomial Regression — Assignment Report

Roll no. IMT2024067 · Code: PyTorch

GitHub: https://github.com/VishudhaSood/Assignment1_ML


1 The task

We get two datasets. In each one we must predict a number y from some inputs, using **only polynomial regression**. Each has 1000 training rows (with y) and 1000 test rows (without y). The teacher scores our test predictions with **MSE** (lower is better) and **R^2** (closer to 1 is better). The main decision is the **degree** of the polynomial.

	var1 — steam turbine	var2 — thermal reservoir
Inputs	6 ($x_1 \dots x_6$)	3 (x_1, x_2, x_3)
Max degree allowed	10	20
Input values	between -1 and 1	between -1 and 1
y in training data	from -9.6 to 12.4	from -29.7 to 30.2

What we noticed in the data. No missing values. Many inputs are exactly -1 or $+1$ (the edge of the allowed range): about 30% in var1 and 25% in var2. The var1 *test* file has even more edge points than the training file, so we checked edge points separately (Section 5).

2 Method



2.1 Polynomial features

A polynomial of degree d uses every term whose powers add up to at most d . For example with two inputs and $d = 2$ the terms are: $1, x_1, x_2, x_1^2, x_1x_2, x_2^2$. The model is just a weighted sum of these terms, so it is still **linear regression**, only on more columns. The number of terms grows fast:

Degree	1	3	5	8	10	20
var1 (6 inputs)	7	84	462	3003	8008	—
var2 (3 inputs)	4	20	56	165	286	1771

With only 1000 rows, a high degree has more terms than rows. Plain least squares would then fit the training data perfectly, noise included — this is **overfitting**. A too-low degree misses the real shape — **underfitting**.

Small extra: Legendre polynomials. Instead of raw powers $x, x^2, x^3, \dots$ we build the terms from Legendre polynomials $P_1(x), P_2(x), P_3(x), \dots$ (e.g. $P_2 = (3x^2 - 1)/2$). They describe **exactly the same set of polynomials**, so the model is unchanged. We use them because on $[-1, 1]$ high powers like x^{18} and x^{20} look almost identical, which makes the computer's equation solving unstable; Legendre terms do not have this problem.

2.2 Fitting with a penalty: Ridge and Lasso

To stop overfitting we add a penalty on the size of the weights w (regularisation):

Method	What we minimise	Effect
Ridge (L2)	error $+\lambda \sum w^2$	shrinks all weights, keeps every term
Lasso (L1)	error $+\alpha \sum w $	sets useless weights to exactly 0 $\rightarrow$ keeps only some terms
Relaxed Lasso	Lasso first, then refit	Lasso chooses the terms; we then refit only those terms with (almost) no penalty, so they are not shrunk too much

λ and α control how strong the penalty is. Ridge has a direct formula,

$$w = (\Phi^\top \Phi + \lambda n I)^{-1} \Phi^\top y,$$

solved with `torch.linalg.solve`. Lasso has no formula, so it is solved step by step (gradient steps + pushing small weights to 0). The constant term (intercept) is never penalised.

2.3 Choosing degree and penalty: 5-fold cross-validation

We split the training rows into 5 equal parts. We train on 4 parts and measure the MSE on the 5th part, which the model has not seen, and repeat so every part is the ‘unseen’ part once. The average is the **CV MSE**, our estimate of the test error. We computed it for **every degree** (1–10 for var1, 1–20 for var2), each method, and several penalty strengths, and picked the lowest. That model was then retrained on all 1000 rows to predict the test file.

3 var1 — steam turbine (6 inputs)




Figure 1: CV MSE vs degree (best penalty for each method). Lower is better.

Degree	No. of terms	Ridge	Lasso	Relaxed Lasso
2	28	2.954	2.947	2.950
3	84	0.959	0.936	0.938
4	210	0.682	0.598	0.610
5	462	0.554	0.358	0.353
6	924	0.801	0.374	0.370
7	1716	1.025	0.383	0.393
8	3003	1.230	0.384	0.395
9	5005	1.412	0.379	0.393
10	8008	1.493	0.380	0.393

Table 1: CV MSE for each degree and method.

Chosen: degree 5, relaxed Lasso ($\alpha = 0.01$). CV MSE = 0.353, CV $R^2 = 0.965$. Only 105 of the 462 possible terms are used.

Why degree 5? From degree 1 to 5 the error drops a lot (lower degrees underfit). After 5 it stops improving, so degree 5 is the simplest model that captures the pattern.

Why Lasso and not Ridge? At degree 5 Ridge's error is 0.55 but Lasso's is about 0.35. Ridge keeps all 462 terms, while Lasso throws most of them away, so the real formula only uses a few terms. Ridge also gets much worse at higher degrees (more and more useless terms), while Lasso stays flat because it just removes them.

4 var2 — thermal reservoir (3 inputs)




Figure 2: CV MSE vs degree. The minimum is at degree 8.

Degree	No. of terms	Ridge	Lasso	Relaxed Lasso
4	35	3.710	3.710	3.710
6	84	0.627	0.627	0.630
7	120	0.363	0.365	0.363
8	165	0.301	0.303	0.304
9	220	0.312	0.311	0.318
10	286	0.351	0.319	0.332
12	455	0.533	0.347	0.368
14	680	0.877	0.413	0.441
16	969	1.129	0.432	0.449
18	1330	1.258	0.455	0.501
20	1771	1.438	0.459	0.513

Table 2: CV MSE for selected degrees.

Chosen: degree 8, Ridge ($\lambda = 0.001$). CV MSE = 0.301, CV $R^2 = 0.993$. All 165 terms are used.

Why degree 8 and not 20? Degree 20 is allowed, but the CV error is lowest at 8 and rises after that. Above 8 the extra terms mainly fit noise (overfitting). Degree 20 Ridge is about $5\times$ worse than degree 8.

Why Ridge? At degree 8 all three methods give almost the same error, so there is no sign that terms should be removed. We keep the simplest method, Ridge, which has a direct formula.

5 How good are the predictions?

The real test answers are hidden, so we estimated test performance in three ways:

Check	var1	var2
Hold-out test: train on 80% of rows, test on the other 20% (10 random splits)	MSE 0.352, R^2 0.964	MSE 0.296, R^2 0.993
Repeated 5-fold CV (3 different shuffles)	MSE 0.353, R^2 0.965	MSE 0.303, R^2 0.993
Test-like MSE: CV error re-weighted so the share of edge points matches the test file	MSE $\approx$ 0.38	MSE $\approx$ 0.31
Baseline: always predict the average y	MSE 9.99, R^2 0	MSE 42.1, R^2 0

All estimates agree, so the results are stable. We expect a test R^2 of about **0.96** for var1 and **0.99** for var2. The var1 test MSE may be slightly higher than the CV value (≈ 0.38) because its test file has more edge points, where errors are a bit larger (0.40 on edge rows vs 0.33 on inner rows).




Figure 3: Predicted vs actual y on unseen CV rows. Points lie close to the diagonal.




Figure 4: Errors (predicted – actual). They are spread randomly around 0 with no pattern, so the model is not missing any systematic shape.

Other checks. Both prediction files have 1000 rows, one column y , no missing values and the same row order as the test files. Re-running the code reproduces them exactly.

6 Summary

	var1	var2
Degree	5	8
Method	Relaxed Lasso, $\alpha = 0.01$	Ridge, $\lambda = 0.001$
Terms used	105 of 462	165 of 165
CV MSE / R^2	0.353 / 0.965	0.301 / 0.993
Hold-out MSE / R^2	0.352 / 0.964	0.296 / 0.993
Prediction file	IMT2024067_pred_var1.csv	IMT2024067_pred_var2.csv

Main takeaways. (1) The best degree is found by cross-validation, not by taking the maximum allowed: 5 out of 10 for var1 and 8 out of 20 for var2. (2) Regularisation is what makes high-degree polynomials usable; Lasso helps most when only a few terms matter (var1). (3) Hold-out and CV estimates agree, so we expect similar errors on the hidden test set.

Code (GitHub: https://github.com/VishudhaSood/Assignment1_ML): `polyreg.py` model, `train.py` degree search + predictions, `validate.py` checks, `make_plots.py`, `make_report.py`.